



UNIVERSITY OF GENEVA
Department of Informatics

Methods and Heuristics for Learning and Optimization

Simulated Annealing and Parallel Tempering for the Traveling
Salesman Problem

14X013

Michel Jean Joseph Donnet

October 25, 2025

Table of Contents

1	Introduction	3
2	Methodology	4
2.1	Modelization and definitions	4
2.1.1	Search space	4
2.1.2	Neighborhood	4
2.1.3	Fitness	5
2.2	Heuristic vs metaheuristic	5
2.2.1	Heuristic	5
2.2.2	Metaheuristic	6
2.2.3	Comparison	6
2.3	Simulated Annealing	7
2.3.1	Physical inspiration	7
2.3.2	Algorithm principe	7
2.3.3	Algorithm structure	7
2.3.4	Algorithm variant: the Parallel Tempering (PT)	9
3	Implementation	11
3.1	Requirements and usage	11
3.2	Pseudo-codes	12
3.2.1	Greedy	12
3.2.2	Simulated annealing	12
3.2.3	Parallel Tempering	13
4	Results	14
4.1	Benchmark problem	14
4.2	Cities problem	16
4.3	Cities2 problem	17
4.4	Random problem	18
4.4.1	20 cities	18
4.4.2	60 cities	20
4.4.3	80 cities	21
4.4.4	100 cities	22
4.5	Enhancement tests	24
5	Discussion	26
5.1	Execution time	26
5.2	Benchmark problem	26
5.3	Cities problems	27
5.3.1	cities.dat	27

5.3.2	cities2.dat	28
5.4	Random problems	28
5.5	Enhancement tests	29
6	Conclusion	30
7	Appendix	32

1 Introduction

(0.25)

be specific on

- search space
- neighborhood
- fitness
- TSP
- greedy
- SA

Human have been always interested in optimization problem to avoid a loss of energy or money or something else. In our digital world, the optimization play an important role in most domains. For instance, a travel compagny will try to maximise its customer base by proposing the customer to visit a maximum number of cities with minimal cost for him. Depending on the optimization of this problem, the compagny will have more customers and will develop more than another compagny in the same domain. To deal with theses optimization issues, heuristics and metaheuristics algorithm are created, using different approaches and modelizations of the real problem to propose an optimization. For instance, in case of the travel compagny, the problem can be modelized by finding a route that minimizes costs and passes through all the cities before returning to the agency. Each city then represents a node in a graph, and each city is connected to the others by an edge weighted by the price of travel between those two cities. This modelization of the problem is called Traveling Salesman Problem (TSP) and is known to be a NP-hard (Nondeterministic Polynomial-time hard) problem, that means that we can't find a solution in polynomial time. To find a solution that may not necessarily be optimal, we use heuristics and metaheuristics algorithm.

In this work, we will try to find a solution to the TSP problem by a metaheuristic method called Simulated Annealing method.

2 Methodology

(1.5)

First, we will explain some terms and definitions used in the rest of this report.

2.1 Modelization and definitions

The TSP problem can be modeled by N cities to visit, numeroted from 0 to $N - 1$, with a cost matrix W that gives at the position i, j the cost of the travel between city i and city j . Each city is connected to each others, so the graph is fully connected. In this work, we assume that the cost of the travel between city i and j is the same in both directions. So the cost matrix W is symetric because $W_{ij} = W_{ji} \forall i, j \in \{0, \dots, N - 1\}$, which implies that $W = W^T$. Each solution can be written as a list of N cities in the order of visit. When the last city of the list is reached, the customers returns to the agency in the first city of the list. For instance, for $N = 3$ cities, a solution can be $[3, 1, 2]$ that means the customer will first visit city 3, then city 1 and finally city 2 and then return to agency in city 3.

2.1.1 Search space

The search space gives the set of all possibles solutions for a given problem. In specific case of TSP problem, the search space will be the set of all the cyclic paths that go through all the cities.

As the paths are cyclic, it means that for instance path $[1, 2, 3]$ is the same that $[3, 1, 2]$. For N cities, for a given path, it's possible to perform N cyclic rotations before returning on the given path.

As we have assumed that the cost between city i and j is the same in both directions, it means that for instance path $[1, 2, 3]$ is the same that $[3, 2, 1]$. So for N cities, for a given path with its set of cyclic rotations, if we perform an inversion, we will obtain an inversed path with its corresponding set of cyclic rotations.

So a path with a given cost can have $2N$ variants.

So for N cities, the search space will be of size S given by the formula 1.

$$S = \frac{\text{Number of possible paths}}{\text{Number of variant per path}} = \frac{N!}{2N} = \frac{(N-1)!}{2} \quad (1)$$

So with a great N , the search space becomes astronomic.

Example: $N = 15 \implies S = \frac{(15-1)!}{2} = 4.3 \times 10^{10}$

2.1.2 Neighborhood

The neighborhood returns a set of solutions closest to the given solution by slightly modifying the given solution. It depends on the modelization of the problem. In our case,

each solution is a list of cities in order of visit.

We define the neighborhood like a swap of 2 elements of the current solution that gives another solution. The neighbor solution obtained must differs from the current solution and also from all cyclic rotations and inversions of the current solution.

So for a solution given with N cities, the size of the neighborhood is borned by the number of unique permutations of 2 elements in a list, like the formula 2 show us. The born is the number of permutations of 2 elements in a list because all the permutations will not be taken since certain permutations can result in a cyclic or inverse rotation of the solution.

$$\text{Number of neighbors} \leq \binom{N}{2} = \frac{N!}{2!(N-2)!} = \frac{N(N-1)}{2} \quad (2)$$

2.1.3 Fitness

The fitness, or cost function, returns the total cost of the travel around all the cities, according to the weight matrix W . As explained above, the cost between city i and city j is given by the element i, j of the cost matrix W which is symetric.

So we can define a fitness function as the sum of the cost of all travel between cities in a given solution x , as shown in the formula 3.

$$f(x) = W_{x_N x_1} + \sum_{i=2}^N W_{x_{i-1} x_i} \quad (3)$$

Note that the term $W_{x_N x_1}$ in the formula 3 design the cost of returning to the agency.

2.2 Heuristic vs metaheuristic

2.2.1 Heuristic

An heuristic approach to an optimization problem will try to find a good solution in polynomial time. It's a deterministic approach, which means that with same input, the method will give always the same output: nothing is given to the hasard. However, the heuristic approach is specific to each problem. Due to its deterministic nature, the algorithm will tend to be blocked by local optimal solutions in case of complex problems.

For instance, an heuristic algorithm that find the minimum of a function will pretty work on simple function like squared functions. However, for a complex function that admits a lot of local minimum, the heuristic algorithm will be blocked in a local minimum because for the algorithm, the minimum is found.

A greedy algorithm is a heuristic algorithm which make the best choice at each step. In this way, we hope to reach the global optimum.

2.2.2 Metaheuristic

A metaheuristic approach will try to find a good solution for a given problem by exploring the search space 2.1.1 of the problem. It's a stochastic approach, that means that probabilities are integrated to choose a solution for next step, that allow to don't take always the optimum solution. The stopping criteria is not the found of a good solution, but it is defined by the user through the guided parameters. The guided parameters are customisable by user and indicate in most case how, where and how long the metaheuristic algorithm will search.

Since the algorithm does not stop when it finds a solution and have some non-deterministic mechanism to choose next solution, theses prevent it from being blocked by a local solution, compared to heuristic algorithms.

So metaheuristic algorithms are good for complex problems. However, for simple problems, heuristic algorithm are more efficient because they will not make a big exploration of the search space like a metaheuristic algorithm can make, especially if the guided parameters are not properly set.

Metaheuristic algorithms are relatively simple to implement and don't depend on given problem. The challenge in metaheuristic algorithms result in the choice of guided parameters, which balance exploring the search space and exploiting the regions already visited.

2.2.3 Comparison

The table 1 describe the differences between heuristic and metaheuristic algorithms.

Table 1: Comparison between heuristic and metaheuristic algorithms

Feature	Heuristic	Metaheuristic
Method type	deterministic	often stochastic
Approach	problem-based	general
Implementation	often simple	often simple
Speed	fast	slower
Problem type	well-defined, structured problems	complex or poorly understood problems
Stopping criteria	fixed (e.g. solution found)	user-defined (e.g. time, iterations, convergence)
Local optima	often stuck	can escape
Use case	when problem structure is known	when heuristics are ineffective

In our case, we will try to implement both heuristic and metaheuristic algorithms: a greedy algorithm for heuristic and an algorithm that follows the simulated annealing

method for metaheuristic. Then, we will compare obtained results.

2.3 Simulated Annealing

The simulated annealing is inspired on the physical annealing process in metallurgy.

2.3.1 Physical inspiration

In metallurgy, to improve quality and capabilities of a metal, the metal is heated to high temperature and then cooled slowly. This is based on the fact that every physical system tends toward a state of equilibrium where energy is minimal. So when the metal is heated, the atoms of metal will have enough energy to move inside the metal. Then, they will position themselves in such a way as to reduce energy consumption. Their new positions will be more ordered than before heating the metal. This will create a better structured metal, which will improve quality and capabilities of the metal, like making it less brittle, more flexible or something else.

2.3.2 Algorithm principle

For the simulated annealing method in algorithmics, it will be the same process.

At the beginning of the algorithm, the temperature will be very high, allowing a great exploration of the search space. In fact, with a great temperature, the system will have enough energy to make big jumps accross the search space.

Then, the temperature will decrease, which will force the exploitation. It's due to the fact that with a smaller temperature, the system is no longer able to achieve big jumps accross the search space.

Finally, the algorithm will run until a state of equilibrium is reached, with a temperature more and more cold. The state of equilibrium is in general user defined. In our case, the state of equilibrium is reached when no improvement of fitness is made during last 3 changes of temperature.

2.3.3 Algorithm structure

The algorithm can be represented by the Figure 34 in appendix 7.

It basically works like this:

- initial configuration and temperature. In general, the initial configuration is choosen randomly. The temperature is a guided parameter. In our case, we take 100 random configurations in the entire search space and compute the mean of the absolute value of the fitness difference ΔE of theses points with the initial configuration. Then, the initial temperature T_0 is defined such as $e^{\frac{-\langle \Delta E \rangle}{T_0}} = 0.5$ with 0.5 the acceptance rate of a given configuration of the search space. The global search allows to give a global view of the variance of fitness inside the search space.

- Update configuration according to the metropolis rule given by the formula 4.

$$P = \begin{cases} 1 & \text{if } \Delta E < 0 \\ e^{-\frac{\Delta E}{T}} & \text{otherwise} \end{cases} \quad (4)$$

ΔE is the fitness difference between the new configuration and the current configuration. So the probability P to choose the given element is always 1 if this element gives an improvement to the current fitness, and is of $e^{-\frac{\Delta E}{T}}$ otherwise, as described in the formula 4. In other words, it means that if the new configuration is an improvement of the current configuration, it will be directly accepted. And if the new configuration is not an improvement of the current configuration, it means that the difference is positive and the new configuration will be accepted with a probability that decrease exponentially as the difference grows, as shown on Figure 1.¹




Figure 1: Probability to choose configuration

- If an equilibrium state is reached, update temperature according to the temperature schedule (a guided parameter). In general, the temperature variation is defined as a geometric sequence, that is given in our case by: $T_{k+1} = 0.9 \times T_k$. The equilibrium state is reached when enough configurations have been accepted or when enough

¹For details about impact of temperature on the search, see section 2.3.3

configurations have been visited (configurations accepted + not accepted) with the given temperature. In our case, to reach equilibrium state, $12N$ configurations must be accepted or $100N$ configurations must be visited with N the number of cities.

- Iterate until the halting condition is reached. The halting condition is in general given by no improvement of the fitness during last X temperature changes (in our case, $X = 3$).
- Return the best configuration that gives the best fitness.

As it is easily possible to see, the method depends on a few guided parameters, like the initial temperature and the update of the temperature... This represents a challenge to set them correctly, and the quality of the result depends on this guided parameters. That's why the main difficulty of the simulated annealing method, like other metaheuristics methods, is, among other things, to correctly set the guided parameters.

Impact of the temperature The temperature plays a crucial role in the simulated annealing method. We speak about high and low temperature, but what is the impact of the temperature on the choice of the next configuration ? To answer, we will see 2 cases.

1. *Low temperature.* When the temperature is low, as we can see on the Figure 1, the probability to accept an element that is worse than the current configuration is very low, so we accept only configurations that are at least as good as the current configuration. The current configuration is therefore exploited since we try to find a better configuration around the given configuration.
2. *High temperature.* When the temperature is high, as the Figure 1 show us, the probability to accept worse configurations is high. So we try to explore further the search space to find better configuration by accepting more configurations.

2.3.4 Algorithm variant: the Parallel Tempering (PT)

There exists different variants of the simulated annealing algorithm, such as parallel tempering.

This variant is based on the same principle, but works quite differently, with different guided parameters. The parallel tempering algorithm can be represented by the diagram 35 in appendix 7.

The principle is to explore space in several locations. To achieve this goal, M replicas of the simulated annealing method are made. A list of M temperature is set such as $T_1 < \dots < T_M$. This list describe the fixed temperature T_i for the simulated annealing replica i . A swap between adjacent replicas is made at a given frequency.

In more details, the parallel tempering works as follows:

- Initial configuration: M random configurations are choosen. In general, $M = \sqrt{N}$.
- Initial temperature: a list of temperature is defined with $T_1 < T_2 < \dots < T_M$.

- For each current configuration i , choose a new configuration according to the Metropolis rule 4 with T_i as temperature.
- Each n iterations, swap configurations according to the fitness of these configurations. The swaps are only attempted between configurations with adjacent temperatures. So for instance, configuration i cannot be swapped with configuration j if there exists a configuration k such as $T_i < T_k < T_j$. Δ_{ij} is then defined as $(fitness_j - fitness_i)(\frac{1}{T_j} - \frac{1}{T_i})$ with $i < j$. The Metropolis rule is then applied to Δ_{ij} to decide if the swap is accepted or not, as described in the formula 5.

$$P = \begin{cases} 1 & \text{if } \Delta_{ij} < 0 \\ e^{-\Delta_{ij}} & \text{otherwise} \end{cases} \quad (5)$$

In this way, as explained in section [-sec-temp], when we have a good configuration, we exploit it by decreasing the temperature, and when we have a bad configuration, we want to explore more the search space to find better configurations by increasing the temperature.

- Iterate until halting condition is reached. In our case, the halting condition is the number of iterations.

So this algorithm can be an improvement of the simulated annealing system depending on the problem, and a parallelization of the method, as its name suggests, is possible to achieve, allowing more exploration of the search space. However, like for each meta-heuristics, the guided parameters such as M , the list of temperatures $T_1 < \dots < T_M$ and the maximum number of iterations must be properly defined to have a method that works quite well.

3 Implementation

The implementations of both simulated annealing and parallel tempering are based on the diagrams 34 and 35, included in the appendix 7. However, some pseudo-code are provided in sections below.

3.1 Requirements and usage

The code is written in python.

The figures are made with the python library Plotly that allows more control on graphs than `matplotlib.pyplot` library.

The requirements are:

- `python 3`
- `numpy`: library used to work with arrays
- `argparse`: library used to manage command line arguments
- `plotly`: python library used to plot the graphics. Documentation can be found on official website: <https://plotly.com/python/>
- `os`: library that allows to interact with operating system
- `time`: library used to capture execution time of algorithms

The code works in the command line, according to the following documentation:

```
usage: tp2.py [-h] [-b BENCHMARK] [-r RANDOM] [-d DATAFILE] [-t THRESHOLD]

TP3 - Traveling Salesman Problem

options:
  -h, --help                show this help message and exit
  -b, --benchmark BENCHMARK
                           Benchmark problem of N cities (trivial solution
                           for humans !)
  -r, --random RANDOM       Random problem of N cities
  -d, --datafile DATAFILE
                           Path to data file for problem with given data
  -t, --threshold THRESHOLD
                           Acceptance rate of algorithms. Default is 50%
```

The output of the program is some graphics and tables of TSP problem solutions for greedy, simulated annealing and parallel temperin methods.

For instance, to generate benchmark graphs for 30 cities, use the following command:

```
python tp2.py -b 30 -t 0.05
```

And to generate graphs for `cities.dat` specific problem, use the following command:

```
python tp2.py -d /path/to/cities.dat
```

3.2 Pseudo-codes

We are looking for path that goes across N cities. The set of cities is given, and the cost of a travel between 2 cities is given by distance between theses 2 cities.

3.2.1 Greedy

- initial step: choose randomly an initial city and put it in the path
- loop until each city is in the path:
 - compute distance between given city and the rest of the cities
 - choose city with smallest distance
 - add city to the path
- return the path

3.2.2 Simulated annealing

- initial step:
 - choose randomly a path. If `enhanced_mode` is activate, choose a path according to an optimization method (combination of greedy and simulated annealing)
 - choose initial temperature:
 - * compute average of fitness change $\langle |\Delta E| \rangle$ for 100 random paths
 - * take initial temperature T_0 such as $e^{-\frac{\langle |\Delta E| \rangle}{T_0}} = threshold$ with *threshold* the acceptance rate of a path in the search space.
- loop until no improvement of fitness is made during last 3 temperature changes
 - choose randomly a new path from current path (by swapping 2 cities)
 - compute probability of acceptance according to Metropolis rule 4
 - if accepted
 - * update path
 - * if updated path is better than best path found so far
 - update best path
 - indicate that an improvement of fitness is made

- if number of paths attempted is greater than $100N$ or number of paths accepted is greater than $12N$
 - * update temperature according to `T_schedule` given: $T_{k+1} = T_schedule \times T_k$
- return best path with best fitness

3.2.3 Parallel Tempering

The number of replicas of simulated annealing is given by M that is generally defined like this: $M = \sqrt{N}$

1. initial step

choose randomly M paths. If `enhanced_mode` is activate, choose M paths according to an optimization method.

choose initial temperature:

- compute average of fitness change $\langle |\Delta E| \rangle$ for 100 random paths
- take initial temperature T_0 such as $e^{-\frac{\langle |\Delta E| \rangle}{T_0}} = threshold$ with *threshold* the acceptance rate of a path in the search space.
- Get temperature list T multiplying T_0 by a sequence of factors that follow a `linear`, `quadratic`, `exponential` or `logarithmic` progression. For instance, for linear progression, we have $T = \frac{1}{N}T_0 < \frac{2}{N}T_0 < \dots < T_0$

2. loop until max number of iterations is reached

for each path in paths

- choose randomly a new path (by swapping 2 cities)
- compute probability of acceptance according to Metropolis rule 4
- if accepted
 - update path
 - if updated path is better than best path found so far
 - * update best path

each `swap_frequency` iteration

- swap adjacent replicas according to their fitness and Metropolis rule 5. For instance, if $fitness_1 < fitness_2$, swap replicas 1 and 2.

3. return best path with best fitness

4 Results

To compute the results, each method is runned 10 times.

The figure with the cities shows the best path during the 10 runs of the algorithm, with on x label the fitness of the given path.

The table give always the average values obtained.

Empirically, experiments have shown that the sequence of factors that provide best results for parallel tempering is in general the quadratic sequence. However, some other sequences progressions proved to be effective. So to compute parallel tempering, the quadratic factor sequence is the sequence choosen.

The maximum number of iterations of parallel tempering has been fixed at 20000 for each experiment.

The initial acceptance rate has been set to 5%. In fact, due to the paths taken randomly on the entire search space, the initial temperature give an overview of the possible difference of fitness of the problem. This difference is quite big, so empirically, to have good results (with a temperature not too big, allowing a good convergence of the method around a solution), the initial acceptance rate has been set to 5%.

The figure of the fitness variation is obtained by taking the best result and detailing the variations in fitness of this result according to iterations.

Algorithms have been runned on benchmark problem with 30 cities, on `cities.dat` and `cities2.dat`, and on random problem generated with 20, 60, 80, 100 cities.

4.1 Benchmark problem

Benchmark: greedy for 30 cities

Attribute	Value
Best fitness	6.2717
Execution time (s)	0.0006




Figure 2: Benchmark problem of 30 cities for the greedy algorithm

Benchmark: simulated annealing for 30 cities

Attribute	Value
Best fitness	9.5473
Number iteration	37257.1
Initial temperature	1.1477
Enhanced mode	false
Initial acceptance rate	5%
Temperature schedule	$T_{k+1} = 0.9 T_k$
Execution time (s)	1.4607

Attribute	Greedy	PT
Execution time ratio	0.0004	4.4143
Fitness ratio	0.6569	1.4778




Figure 3: Benchmark problem of 30 cities for the simulated annealing algorithm

Benchmark: parallel tempering for 30 cities

Attribute	Value
Best fitness	14.1086
Number iteration	6229.8
Initial temperature	1.0216
Temperature schedule type	quad
Enhanced mode	false
Initial acceptance rate	5%
Execution time (s)	6.4481

Attribute	Greedy	SA
Execution time ratio	0.0001	0.2265
Fitness ratio	0.4445	0.6767




Figure 4: Benchmark problem of 30 cities for the parallel tempering algorithm





Figure 5: Benchmark problem of 30 cities, fitness variation

4.2 Cities problem

cities.dat: greedy for 18 cities

Attribute	Value
Best fitness	33.0499
Execution time (s)	0.0005




Figure 6: cities.dat problem of 18 cities for the greedy algorithm

cities.dat: simulated annealing for 18 cities

Attribute	Value
Best fitness	34.6358
Number iteration	5194.4
Initial temperature	2.4404
Enhanced mode	false
Initial acceptance rate	5%
Temperature schedule	$T_{k+1} = 0.95 T_k$
Execution time (s)	0.2143

Attribute	Greedy	PT
Execution time ratio	0.0025	25.2853
Fitness ratio	0.9542	0.8428




Figure 7: cities.dat problem of 18 cities for the simulated annealing algorithm

cities.dat: parallel tempering for 18 cities

Attribute	Value
Best fitness	29.1924
Number iteration	9582.8
Initial temperature	2.1207
Temperature schedule type	quad
Enhanced mode	false
Initial acceptance rate	5%
Execution time (s)	5.4187

Attribute	Greedy	SA
Execution time ratio	0.0001	0.0395
Fitness ratio	1.1321	1.1865




Figure 8: cities.dat problem of 18 cities for the parallel tempering algorithm




Figure 9: cities.dat problem of 18 cities, fitness variation

4.3 Cities2 problem

cities2.dat: greedy for 49 cities

Attribute	Value
Best fitness	3.4378
Execution time (s)	0.0008




Figure 10: cities2.dat problem of 49 cities for the greedy algorithm

cities2.dat: simulated annealing for 49 cities

Attribute	Value
Best fitness	4.7105
Number iteration	35815.3
Initial temperature	0.1507
Enhanced mode	false
Initial acceptance rate	5%
Temperature schedule	$T_{k+1} = 0.95 T_k$
Execution time (s)	2.0806

Attribute	Greedy	PT
Execution time ratio	0.0004	3.5193
Fitness ratio	0.7298	0.8992




Figure 11: cities2.dat problem of 49 cities for the simulated annealing algorithm

cities2.dat: parallel tempering for 49 cities

Attribute	Value
Best fitness	4.2354
Number iteration	9987.4
Initial temperature	0.1728
Temperature schedule type	quad
Enhanced mode	false
Initial acceptance rate	5%
Execution time (s)	7.3224

Attribute	Greedy	SA
Execution time ratio	0.0001	0.2841
Fitness ratio	0.8117	1.1122




Figure 12: cities2.dat problem of 49 cities for the parallel tempering algorithm





Figure 13: cities2.dat problem of 49 cities, fitness variation

4.4 Random problem

4.4.1 20 cities

Random: greedy for 20 cities

Attribute	Value
Best fitness	4.6648
Execution time (s)	0.0001




Figure 14: Random problem of 20 cities for the greedy algorithm

Random: simulated annealing for 20 cities

Attribute	Value
Best fitness	4.6323
Number iteration	8801.6
Initial temperature	0.2571
Enhanced mode	false
Initial acceptance rate	5%
Temperature schedule	$T_{i+1} = 0.95 T_i$
Execution time (s)	0.3544

Attribute	Greedy	PT
Execution time ratio	0.0004	14.0859
Fitness ratio	1.007	0.919




Figure 15: Random problem of 20 cities for the simulated annealing algorithm

Random: parallel tempering for 20 cities

Attribute	Value
Best fitness	4.2569
Number iteration	11131.9
Initial temperature	0.3421
Temperature schedule type	quad
Enhanced mode	false
Initial acceptance rate	5%
Execution time (s)	4.9925

Attribute	Greedy	SA
Execution time ratio	0	0.071
Fitness ratio	1.0958	1.0882




Figure 16: Random problem of 20 cities for the parallel tempering algorithm





Figure 17: Random problem of 20 cities, fitness variation

4.4.2 60 cities

Random: greedy for 60 cities

Attribute	Value
Best fitness	7.2433
Execution time (s)	0.0004




Figure 18: Random problem of 60 cities for the greedy algorithm

Random: simulated annealing for 60 cities

Attribute	Value
Best fitness	17.0678
Number iteration	34425.8
Initial temperature	0.5886
Enhanced mode	false
Initial acceptance rate	5%
Temperature schedule	$T_{i+1} = 0.95 T_i$
Execution time (s)	1.836

Attribute	Greedy	PT
Execution time ratio	0.0002	5.3077
Fitness ratio	0.4244	0.8187




Figure 19: Random problem of 60 cities for the simulated annealing algorithm

Random: parallel tempering for 60 cities

Attribute	Value
Best fitness	13.9731
Number iteration	8752.8
Initial temperature	0.5725
Temperature schedule type	quad
Enhanced mode	false
Initial acceptance rate	5%
Execution time (s)	9.7449

Attribute	Greedy	SA
Execution time ratio	0	0.1884
Fitness ratio	0.5184	1.2215




Figure 20: Random problem of 60 cities for the parallel tempering algorithm




Figure 21: Random problem of 60 cities, fitness variation

4.4.3 80 cities

Random: greedy for 80 cities

Attribute	Value
Best fitness	9.3829
Execution time (s)	0.0006




Figure 22: Random problem of 80 cities for the greedy algorithm

Random: simulated annealing for 80 cities

Attribute	Value
Best fitness	23.075
Number iteration	54755.1
Initial temperature	0.6685
Enhanced mode	false
Initial acceptance rate	5%
Temperature schedule	$T_{k+1} = 0.95 T_k$
Execution time (s)	2.4948

Attribute	Greedy	PT
Execution time ratio	0.0002	3.9426
Fitness ratio	0.4066	0.8463




Figure 23: Random problem of 80 cities for the simulated annealing algorithm

Random: parallel tempering for 80 cities

Attribute	Value
Best fitness	19.5278
Number iteration	11551.4
Initial temperature	0.7533
Temperature schedule type	quad
Enhanced mode	false
Initial acceptance rate	5%
Execution time (s)	9.8361

Attribute	Greedy	SA
Execution time ratio	0.0001	0.2536
Fitness ratio	0.4805	1.1816




Figure 24: Random problem of 80 cities for the parallel tempering algorithm





Figure 25: Random problem of 80 cities, fitness variation

4.4.4 100 cities

Random: greedy for 100 cities

Attribute	Value
Best fitness	10.1308
Execution time (s)	0.0007




Figure 26: Random problem of 100 cities for the greedy algorithm

Random: simulated annealing for 100 cities

Attribute	Value
Best fitness	31.6565
Number iteration	59183.6
Initial temperature	0.9531
Enhanced mode	false
Initial acceptance rate	5%
Temperature schedule	$T_{i+1} = 0.95 T_i$
Execution time (s)	3.9886

Attribute	Greedy	PT
Execution time ratio	0.0002	2.8235
Fitness ratio	0.32	0.8259




Figure 27: Random problem of 100 cities for the simulated annealing algorithm

Random: parallel tempering for 100 cities

Attribute	Value
Best fitness	26.1444
Number iteration	11626.3
Initial temperature	0.9708
Temperature schedule type	quad
Enhanced mode	false
Initial acceptance rate	5%
Execution time (s)	11.2618

Attribute	Greedy	SA
Execution time ratio	0.0001	0.3542
Fitness ratio	0.3875	1.2108




Figure 28: Random problem of 100 cities for the parallel tempering algorithm





Figure 29: Random problem of 100 cities, fitness variation

4.5 Enhancement tests

Random: greedy for 100 cities

Attribute	Value
Best fitness	8.7578
Execution time (s)	0.0008




Figure 30: Random problem of 100 cities for the greedy algorithm

Random: simulated annealing for 100 cities

Attribute	Value
Best fitness	8.6144
Number iteration	4882.6
Initial temperature	13.8302
Enhanced mode	true
Initial acceptance rate	5%
Temperature schedule	$T_{i+1} = 0.95 T_i$
Execution time (s)	0.2604

Attribute	Greedy	PT
Execution time ratio	0.0029	39.7915
Fitness ratio	1.0167	0.9786




Figure 31: Enhancement for the simulated annealing algorithm

Random: parallel tempering for 100 cities

Attribute	Value
Best fitness	8.43
Number iteration	0
Initial temperature	13.7957
Temperature schedule type	quad
Enhanced mode	true
Initial acceptance rate	5%
Execution time (s)	10.3633

Attribute	Greedy	SA
Execution time ratio	0.0001	0.0251
Fitness ratio	1.0389	1.0219




Figure 32: Enhancement for the parallel tempering algorithm




Figure 33: Fitness variation of enhanced methods

5 Discussion

In this section, we will discuss the results obtained and displayed in the previous section.

5.1 Execution time

First, what's about the execution time of the three algorithms ?

If we look at the problem from the perspective of computation time, the greedy algorithm is $2500\times$ more fast than the simulated annealing and $10000\times$ more fast than the parallel tempering for the benchmark problem for instance, as shown on figures 3 and 4. It's logic because the greedy method construct the path whereas the two other algorithms try different paths at each iterations. It means that only a computation is necessary for greedy algorithm whereas for the other algorithms, a number of computation at least as big as the number of iterations is necessary. So greedy algorithm will be always faster than the others.

According to the comparison Table 1, it's obviously normal.

The execution time depends among other things on the number of computations to achieve.

So the simulated annealing and the parallel tempering method will have a computation time according to the number of iterations achieved. As the parallel tempering applies the simulated annealing on each of the M replica, for a same number of iterations, the method will be slower than the simulated annealing. And suppose that the algorithm is perfectly parallelized, it will always be slower for the same number of iterations because at given frequency, the process units must pool their work to make the swaps between replicas, what takes time.

So the parallel tempering algorithm will be always slower than the simulated annealing for a same number of iterations.

However the parallel tempering could prove more effective than simulated annealing.

That's why we will try to see advantages and disadvantages of both methods for different problems.

5.2 Benchmark problem

The benchmark problem defined is N cities distributed equally around a circle of radius 1. In this way, it's easy to know the optimal fitness: it's the path that create the circle and that have a fitness of $2\pi \approx 6.2$.

When we look at the results, we can see that the greedy algorithm find the optimal solution for this benchmark problem, as shown on Figure 2. It's the expected result, so the algorithm is well implemented.

However, we can constate that on the 10 executions, the simulated annealing doesn't find always the optimal result because the best fitness mean exceeds that of the greedy

best fitness mean, as illustrated in Figure 3.

And for parallel tempering, the optimal solution is never reached, even with 20000 iterations.

So at first glance, the simulated annealing appears better than the parallel tempering.

If we look at the Figure 5 that shows the fitness variation depending on iterations, we can see that for the simulated annealing, the fitness start high (around 40). Then, it decrease, with a variation of fitness more and more small, until reaching a solution. The solution reached is a good solution, but not the optimal because the solution from parallel tempering method is better.

And if we look at the parallel tempering plot, there is a lot of curves. Each curve correspond to a fixed temperature, with T_0 the lowest temperature. It's possible to see that more high is the temperature, more high is the fitness. It's the expected result, because we switch replicas at a given frequency to cool down replicas that have better results to favorise exploitation and we heat up replicas that have worst results to favorise exploration. If we look at T_0 , we can constate that there is only a curve until 2000 iterations whereas the algorithm runs on 10000 iterations. It's normal because in the implementation, the fitness is captured only when the move is accepted according to Metropolis rule 4. And with T_0 , a colder temperature, the Metropolis rule will not accept changes that are worst compared to actual path, as shown on the precedent 1. So less changes are accepted than for a higher temperature, so the curve goes only until 2000 because for this temperature, only 2000 moves have been accepted.

5.3 Cities problems

5.3.1 cities.dat

On this problem, all methods have been runned. If we look at the results (Figure 6, Figure 7 and Figure 8), we can see that the mean of the best fitness is the greatest for simulated annealing, and the slowest for parallel tempering.

So it means that the parallel tempering method has proven to be more effective on this problem. However, all problems have given a good solution because the ratio of the fitness is closest to 1, as shown on Figure 7 and Figure 8.

So it basically means that both methods are approximatively equivalent for this given problem in term of solution. But in term of computation, the greedy is always the fastest, and the simulated annealing is 25× more fast than the parallel tempering, with only a mean of 5000 iterations against 10000 iterations for parallel tempering. It could signify that with more iterations on the simulated annealing method, the results could have been better.

And to have more iterations, the temperature schedule or the initial acceptance rate must be changed.

If we take a look on the fitness variation graph 9, we can see that the simulated

annealing method varies a lot and does not stabilize. So perhaps for this problem, the guided parameters of the method were not properly initialized, that explain the variance of results because of no good convergence to a solution. Or perhaps it's due to stochasticity of the algorithm: it stumbled upon a good solution, making no fitness improvement during temperature changes, which satisfies the halting condition. The number of iterations for the best result is only around 2000 iterations whereas the average of iterations is around 9600, that's why these variations are probably more due to stochasticity of the algorithm than to the bad values of guided parameters.

5.3.2 cities2.dat

On this problem, according to the figures 10, 11 and 12, the greedy algorithm has the best average fitness but the simulated annealing has found the best solution compared to other algorithms. And this time, the time execution of parallel tempering and simulated annealing is approximatively the same, because this time, simulated annealing is only $3.5\times$ more fast than parallel tempering. If we look at the fitness average of simulated annealing and parallel tempering, the parallel tempering has always a better fitness average. However, the best solution of the simulated annealing method is better. In fact, taking a look at the average fitness and the best fitness of boths methods allows us to see that the parallel tempering give always results that have approximatively the same fitness because the difference between the best fitness and the average fitness is around 0.2 whereas for simulated annealing, this difference is arround 2. It basically means that the parallel tempering method give perhaps worst results than simulated annealing, but always gives approximately the same results whereas simulated annealing gives sometimes very good results and sometimes very bad results.

The Figure 13 shows us that for simulated annealing, the guided parameters was correctly set to reach the solution because the curve decrease more and more until it stabilize around optimal solution. In fact, the curve decrease according to the temperature schedule: the temperature is cooled down as time goes by, promoting exploitation of solutions more than exploration. It's the expected behaviour and the solution obtained is very good, so for this problem and this execution, the simulated annealing worked quite well. However, for another execution with the same parameters, the simulated annealing method worked badly because the average best fitness is worst than the average best fitness of parallel tempering.

So simulated annealing can give good results, but parallel tempering give constant results.

5.4 Random problems

On the random problem with 20 cities, the SA algorithm gives the best results. However, for all other random problems, we can see that the greedy algorithm is far superior to the

others. For example, the greedy algorithm is on average $3\times$ better in terms of average best fitness for the problem with 100 cities, as shown in Figure 26, Figure 27 and Figure 28.

If we look at the fitness average between simulated annealing and parallel tempering, once again, we can see that the variance of the fitness is greater in the case of simulated annealing compared to parallel tempering. This result means that in general, results given by parallel tempering are more constant than results given by simulated annealing. And the average best fitness of the parallel tempering algorithm is in general smaller than average best fitness of simulated annealing, like for 60, 80 and 100 cities.

So the simulated annealing method gives better results, but is generally less accurate than the parallel tempering method for this kind of problems.

To conclude, as optimization method was almost always better than metaheuristics methods, it is preferable to use an optimization method for a problem if any optimization method exist because results are as good as metaheuristics and computation time is by far better than metaheuristics algorithms.

5.5 Enhancement tests

To increase accuracy of methods, heuristic and metaheuristic have been combined. The idea is to use an optimization method, like greedy algorithm, to give an initial configuration to the metaheuristic simulated annealing or parallel tempering methods. It's possible to make this because as we have seen, the greedy algorithm get fastly a good solution. So it allows to guide the search in the search space directly around a good solution. And with the stochasticity of the metaheuristic method, we can escape the potential blockage of the heuristic method in a local solution.

The results obtained (Figure 30, Figure 31 and Figure 32) are globally better than greedy algorithm because average best fitness is smaller and optimal solution has also a smaller fitness than the greedy algorithm.

However, if we look at the Figure 33, we can see that for instance, for the simulated annealing method, we begin at low fitness and then the method don't find any enhancement to the greedy fitness. And it's the same for parallel tempering.

So the metaheuristics methods don't find a better solution for this problem by starting with a good solution. And all the search made to explore the search space was costly in computation time (At least $100\times$ slower than greedy algorithm...) So perhaps in this case, it would have been wiser to execute 100 times more greedy algorithm instead of executing a metaheuristic algorithm that is costly in computation time, or to better set the guided parameters.

6 Conclusion

The Traveling Salesman Problem is a NP-Hard problem for which it's difficult to find optimal solutions due to the size of the search space. To find solutions, greedy, simulated annealing and parallel tempering methods were implemented. The greedy method is an efficient way to have quickly a good solution for the given problem. The simulated annealing method, inspired from physical process, is a metaheuristic method that gives good results in term of problem given. Indeed, when tested on a given problem, this method yielded the best results for big problems. The parallel tempering method is a metaheuristic method based on multiple instance of simulated annealing methods working together. Compared to the simulated annealing method, results obtained are more constants because in general, the average of the quality of obtained solution was better than simulated annealing, and sometime than greedy.

Some improvement of both simulated annealing and parallel tempering methods were tested, thus reaching the conclusion that perhaps for the given problem, running several times greedy algorithm could be more promising.

To conclude, results of metaheuristics methods have not performed very well compared to heuristics methods. That's why if an optimization method exist for the given problem, it's more interesting to use it instead of a metaheuristic method that performs an intelligent but exhaustive exploration of the search space.

The table Table 2 summarizes the observations briefly.

Table 2: Comparison between heuristic and metaheuristics methods for Traveling Salesman Problem

	Greedy	Simulated Annealing	Parallel tempering
Execution time	very fast	slow	very slow
Quality of solution	very good	good	medium
Stability of solutions	medium	medium	good
Guided parameters	None	initial acceptance rate, temperature schedule	initial acceptance rate, temperature progression, number of replicas
Enhanced mode		good	good
Usage	when possible	when heuristic not possible	when heuristic not possible

We may wonder if there exists a more efficient way to properly define guided parameters in metaheuristic methods instead of empirically testing values for each of them.

7 Appendix




Figure 34: Simulated Annealing Diagram




Figure 35: Parallel Tempering Diagram